

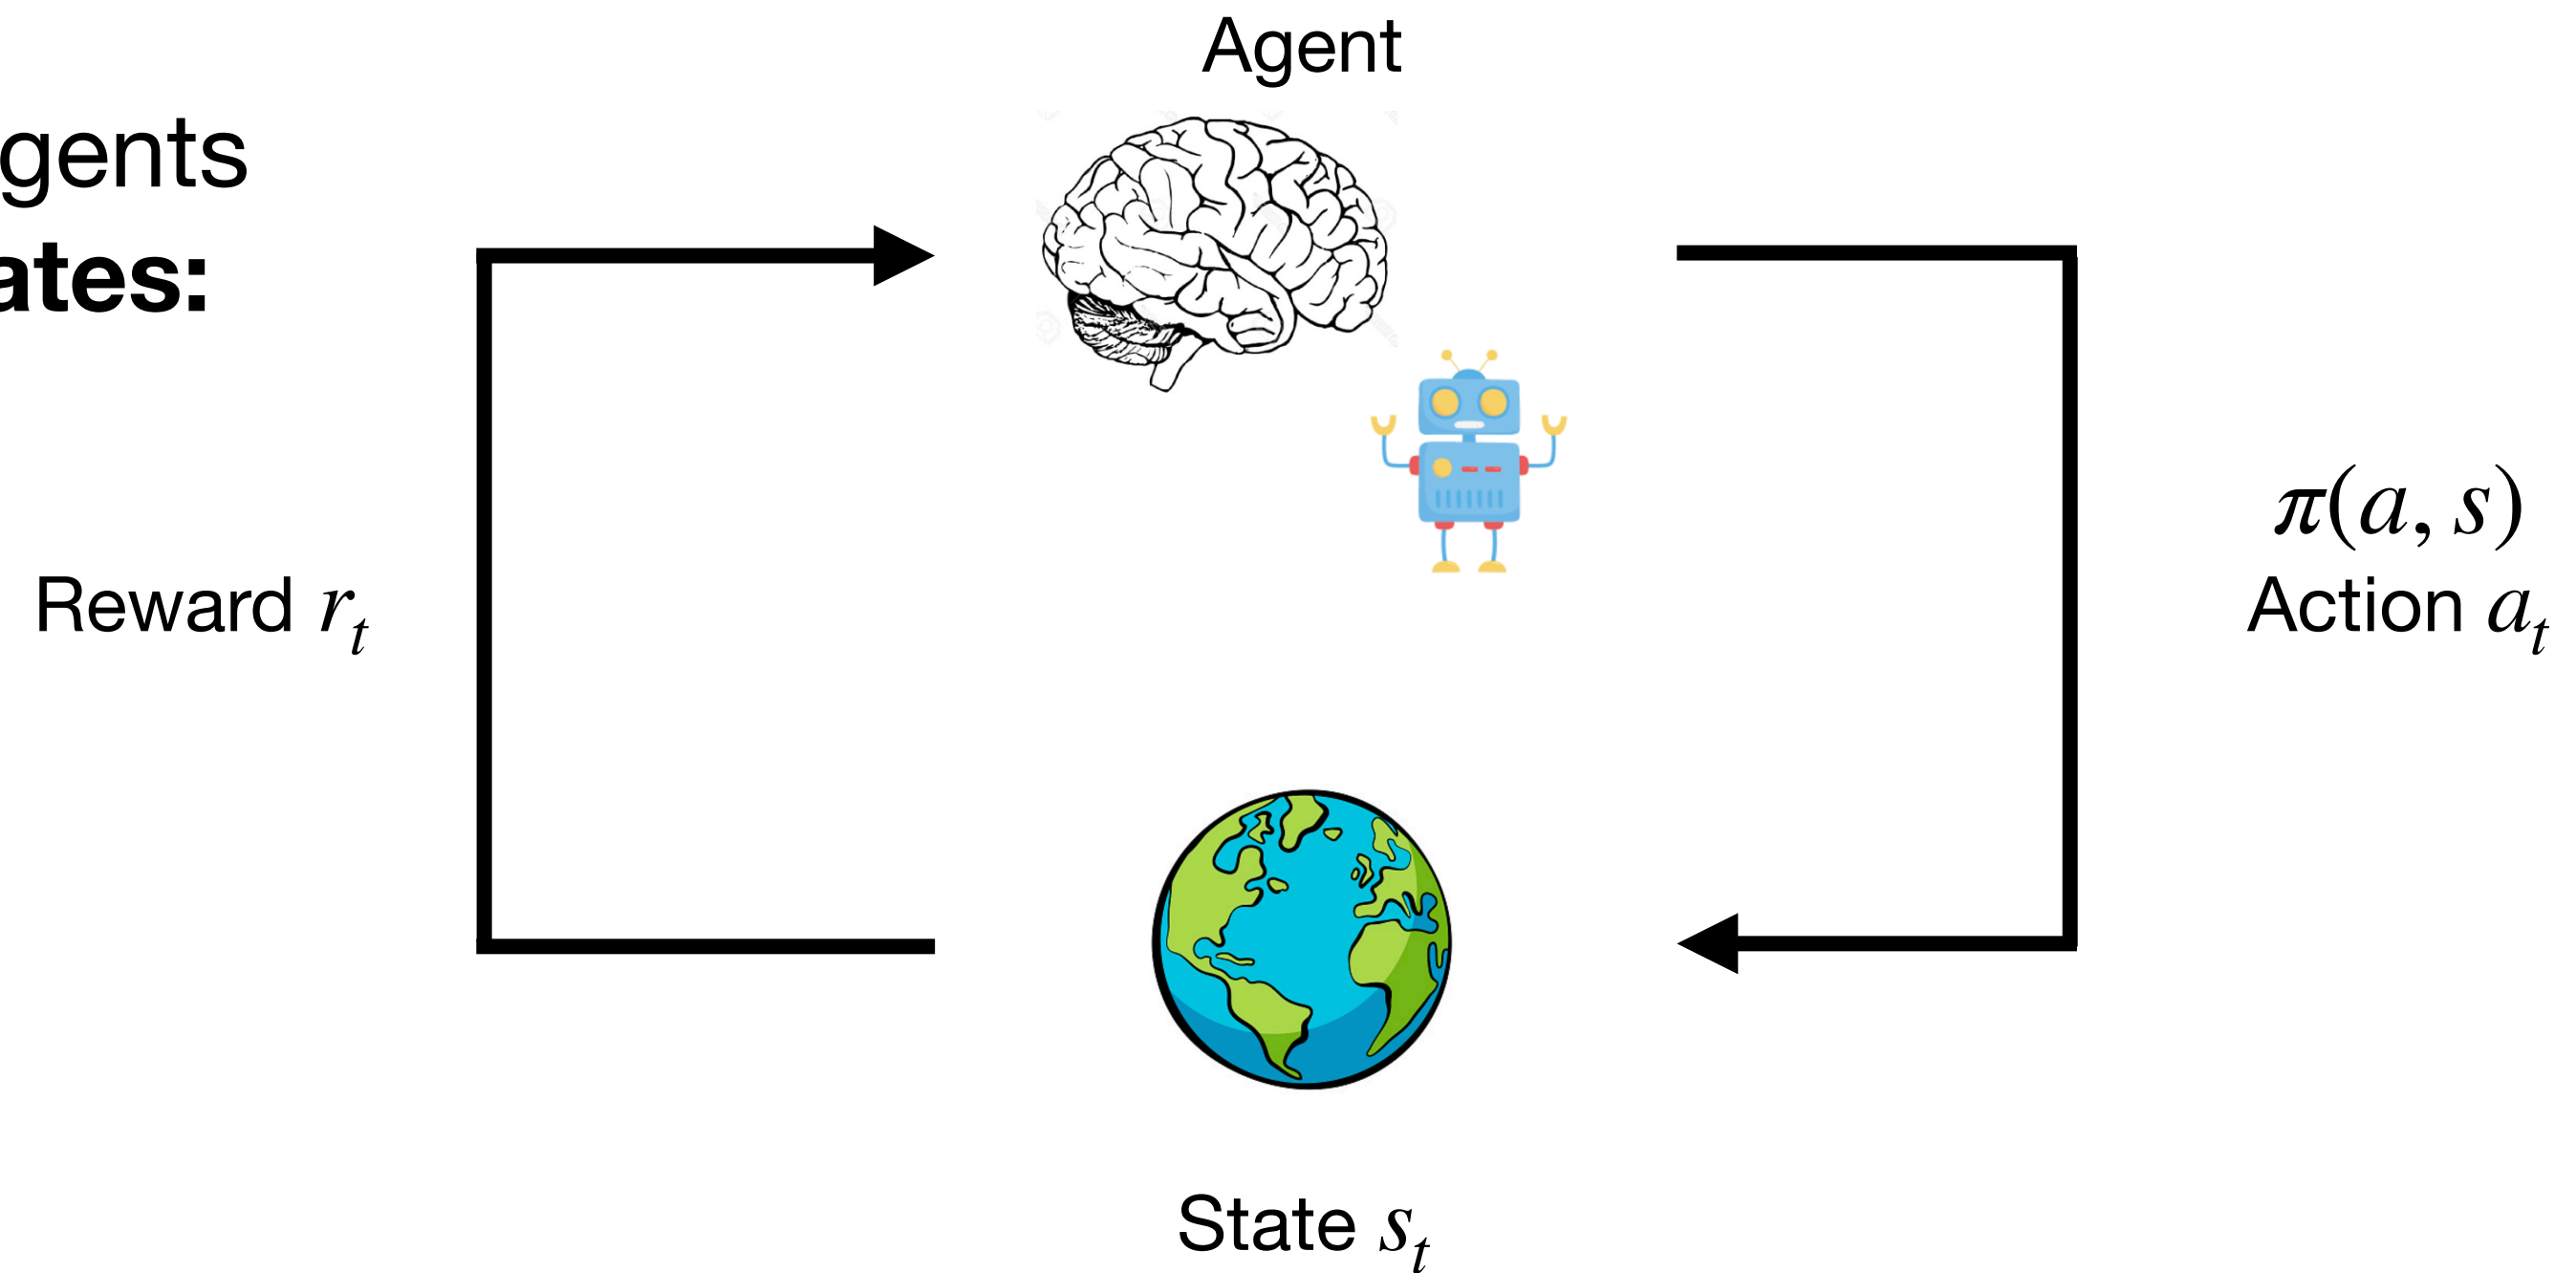
An introduction to Reinforcement Learning

31st of May 2022

Recap: Temporal Difference Learning

Based on a reward signal, agents learn **values of actions/states**:

$$V_{\pi}(s) = \mathbb{E}_{\pi}[R \mid s_0 = s]$$



TD Learning:

$$V(s_t) \leftarrow V(s_t) + \alpha \cdot (r + \gamma \cdot V(s_{t+1}) - V(s_t))$$

Prediction error

Learning rate

Discount rate

Rescorla Wagner Learning:

$$V(s_t) \leftarrow V(s_t) + \alpha \cdot (r - V(s_t))$$



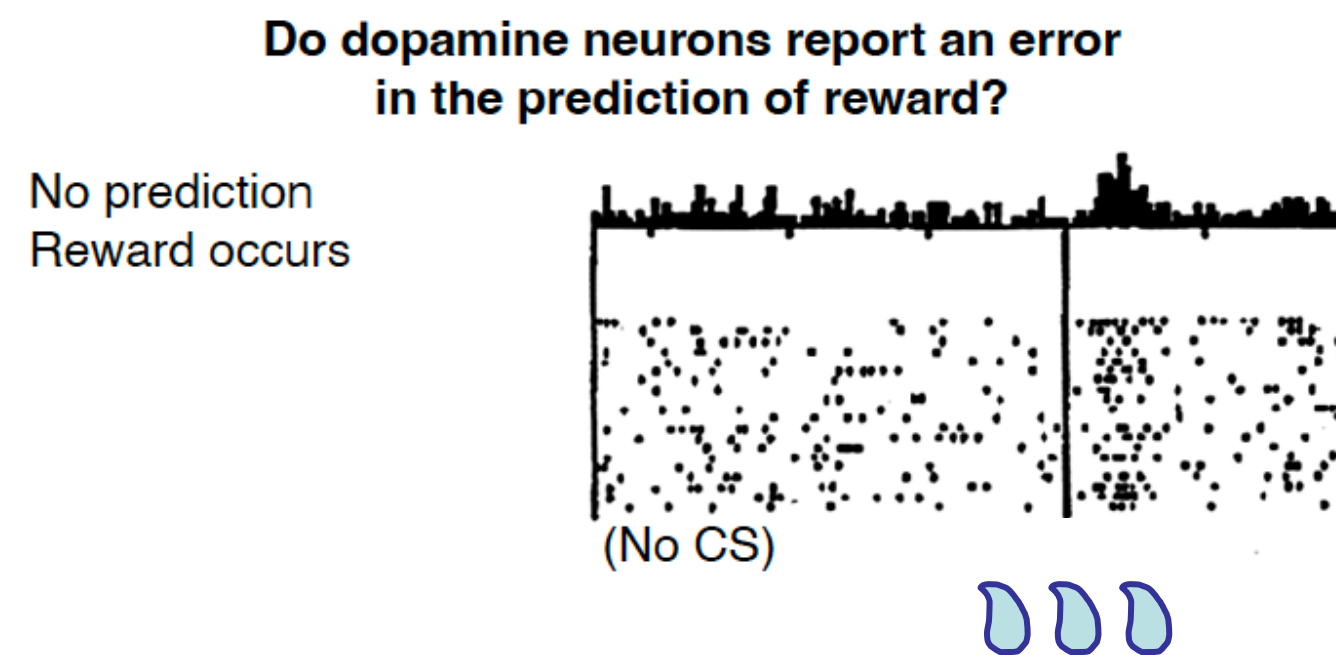

Prediction error

Learning rate

Recap: Can RL tell us anything about the brain?

$$V(s_t) \leftarrow V(s_t) + \alpha \cdot (r + \gamma \cdot V(s_{t+1}) - V(s_t))$$

- It looks like DA signals the reward prediction error in TD learning (Schultz, Dayan & Montague Science, 1997)



Coding: TD Learning

https://github.com/schwartenbeckph/RL-Course/tree/main/2022_05_24