

# **An introduction to Reinforcement Learning**

**26th of July 2022**

# Summary

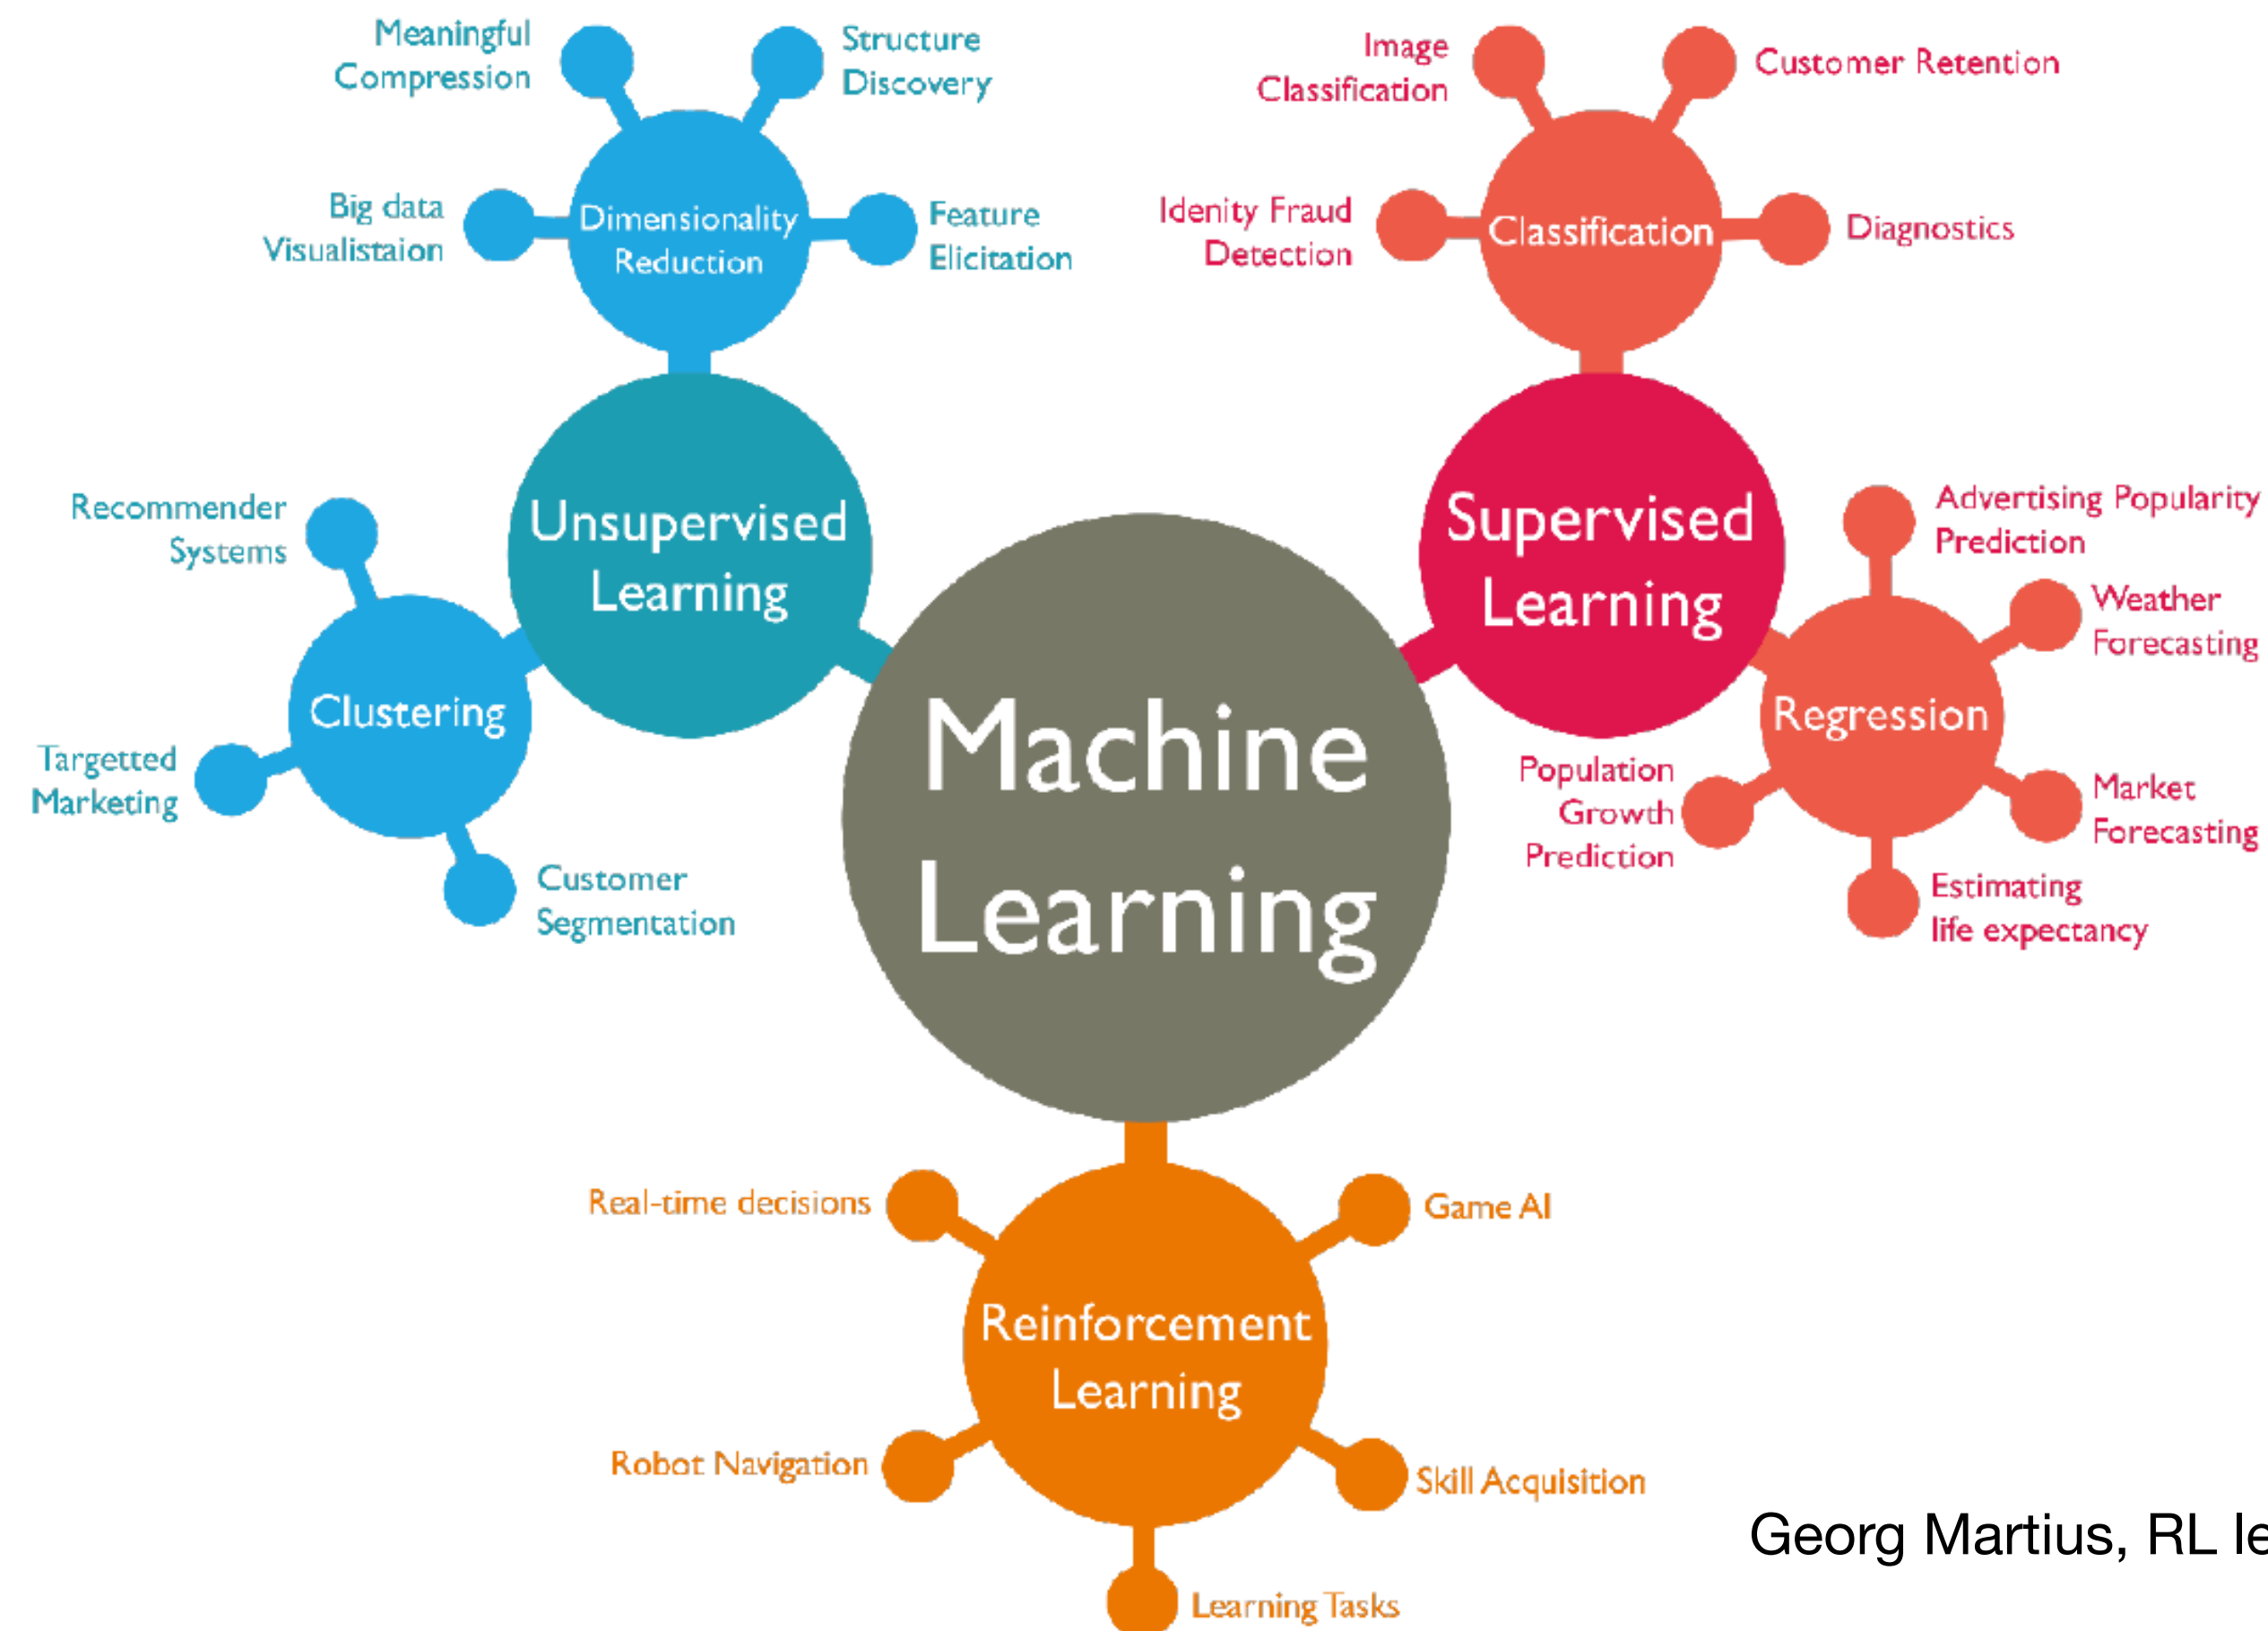
# What is reinforcement learning (RL)?

- RL is a **computational approach** to learning from **interactions** with the **environment**
  - Trial-and-error
  - Delayed reward
- Considers whole problem of **goal-directed** agent interacting with an **uncertain** environment
- RL agents
  - Have explicit goals
  - Sense aspects of their environments
  - Choose actions to influence their environments
- Very general

# How does RL compare to other types of learning?

- Association Learning
  - Representation Learning
  - Supervised Learning
  - Unsupervised Learning
  - Imitation Learning
  - ...
- Reinforcement Learning

# The Machine Learning view:



Georg Martius, RL lecture 2020

# Types of (machine) learning: supervised learning


- Find correct labelling/prediction of data:

Classification



MNIST

Regression





Bishop 2006

- That's not what we want though:
  - Want to learn from *own experience* by *interacting* with the world

# Types of (machine) learning: unsupervised learning

- Find structure in data:



- That's also not what we want:
  - Don't (necessarily) want to learn hidden structure, rather: maximise reward


# Basic setup: how to agents learn to act?

Based on a reward signal, agents learn **values of actions/states**:

$$V_{\pi}(s) = \mathbb{E}_{\pi}[R | s_0 = s]$$

Action is governed by a **policy**:

$$\pi(a, s) = P(a_t = a | s_t = s)$$



Agents can learn a **model of the environment** to make smarter decisions, e.g.:


$$P(s_{t+1} = s | s_t = s, a_t = a)$$



# Basic setup: how to agents learn to act?

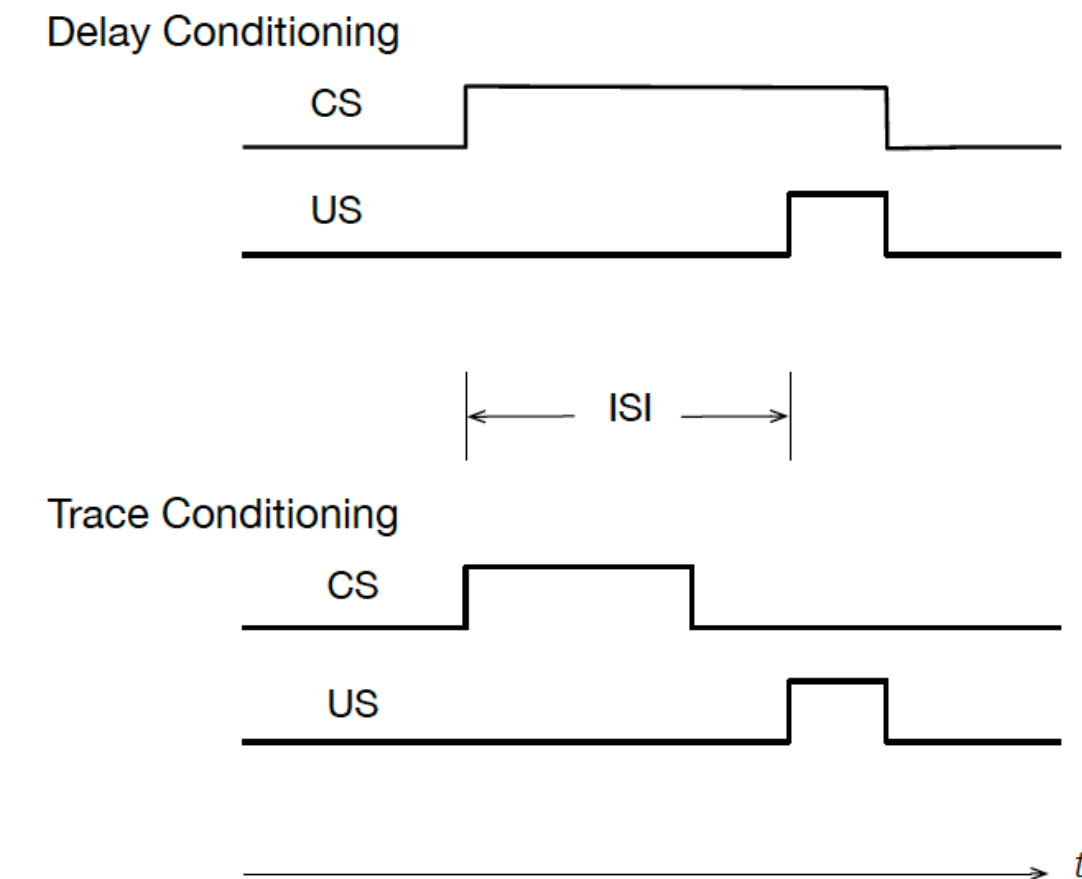
Based on a reward signal, agents learn **values of actions/states**:

$$V_{\pi}(s) = \mathbb{E}_{\pi}[R | s_0 = s]$$




# Learning to predict reward

- Two simple learning algorithms:



- Rescorla-Wagner Learning**


E.g.: blocking



$$V(s_t) \leftarrow V(s_t) + \alpha \cdot (r - V(s_t))$$

Learning rate

- Temporal-Difference Learning** E.g.: Higher-order conditioning



$$V(s_t) \leftarrow V(s_t) + \alpha \cdot (r + \gamma \cdot V(s_{t+1}) - V(s_t))$$

# Recap: Temporal Difference Learning

**TD Learning:** 
$$V(s_t) \leftarrow V(s_t) + \alpha \cdot (r + \gamma \cdot V(s_{t+1}) - V(s_t))$$

Prediction error

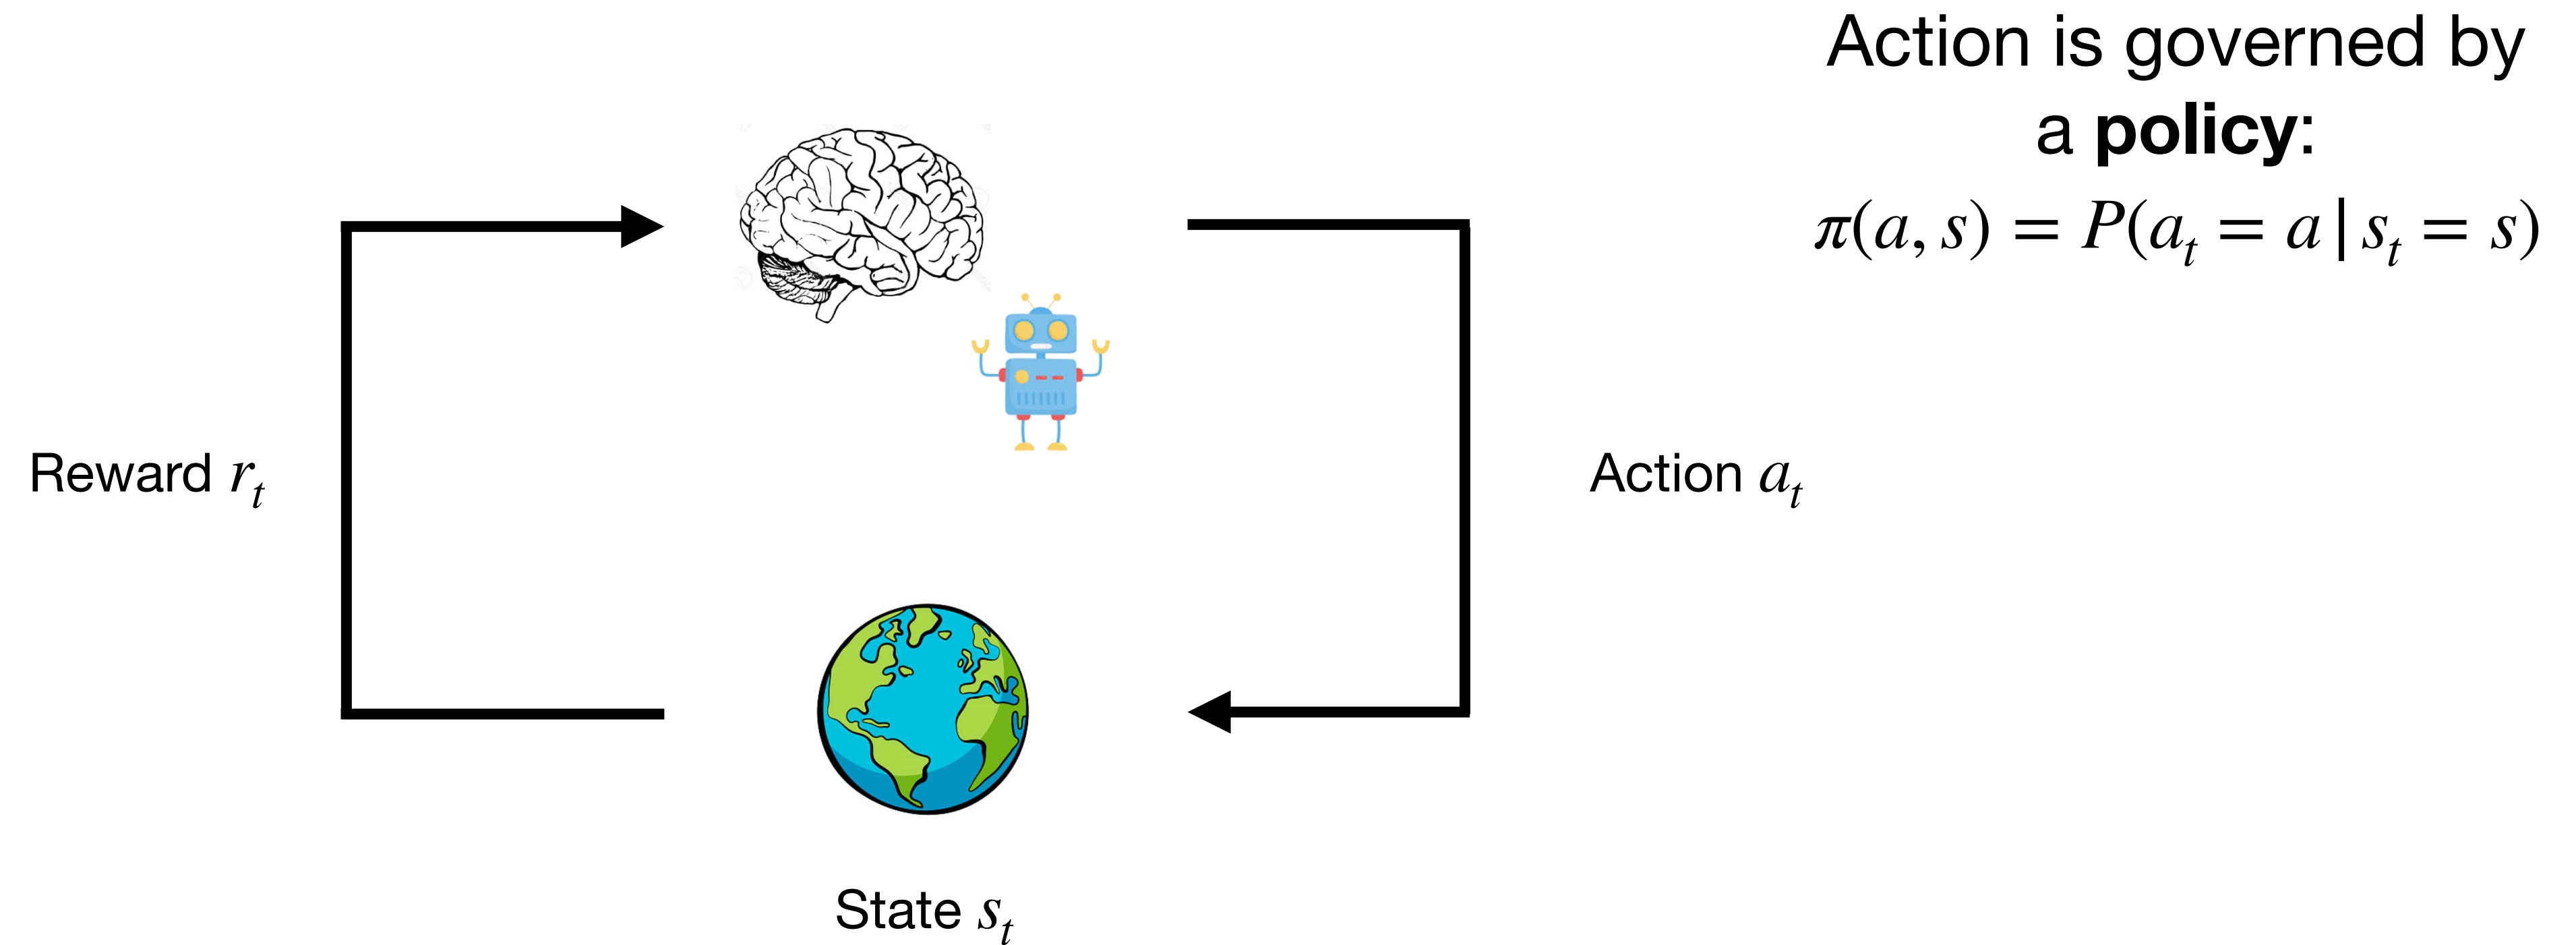
Learning rate

Discount rate



But what about actions?


# Basic setup: how to agents learn to act?



# Multi-armed bandits

**Greedy** action selection:

$$P(a_t = a) = \begin{cases} 1 & \text{if } a_t = \operatorname{argmax}_a V_t(a) \\ 0 & \text{otherwise} \end{cases}$$



**Epsilon-greedy** action selection:

$$P(a_t = a) = \begin{cases} 1 - \epsilon & \text{if } a_t = \operatorname{argmax}_a V_t(a) \\ \epsilon/N & \text{otherwise} \end{cases}$$

**Softmax** action selection:

$$P(a_t = a) = \frac{e^{V_t(a) \cdot \beta}}{\sum_{i=1}^N e^{V_t(a_i) \cdot \beta}}$$

Action is governed by a **policy**:

$$\pi(a, s) = P(a_t = a \mid s_t = s)$$

**Upper-confidence-bound (UCB)** action selection:


$$P(a_t = a) = \operatorname{argmax}_a [V_t(a) + c \cdot \sqrt{\frac{\ln t}{N_t(a)}}]$$

# Limitation of multi-armed bandit problems

Your current action does not influence what happens next!!

How can we solve sequential problems?

The textbook problem:  
**‘Cliff-World’**





# From classical to instrumental learning

## TD Learning:


$$V(s_t) \leftarrow V(s_t) + \underset{\text{Learning rate}}{\alpha} \cdot \left( \underset{\text{Discount rate}}{r + \gamma \cdot V(s_{t+1})} - V(s_t) \right)$$

Prediction error


## Q-Learning:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \underset{\text{Learning rate}}{\alpha} \cdot \left( \underset{\text{Discount rate}}{r + \gamma \cdot \max_a Q(s_{t+1}, a_{t+1})} - Q(s_t, a_t) \right)$$

Prediction error



# Basic setup: how do agents learn to act?




Agents can learn a **model of the environment** to make smarter decisions, e.g.:

$$P(s_{t+1} = s \mid s_t = s, a_t = a)$$



# Markov Process

Most RL problems are problems where agents face sequences of states:

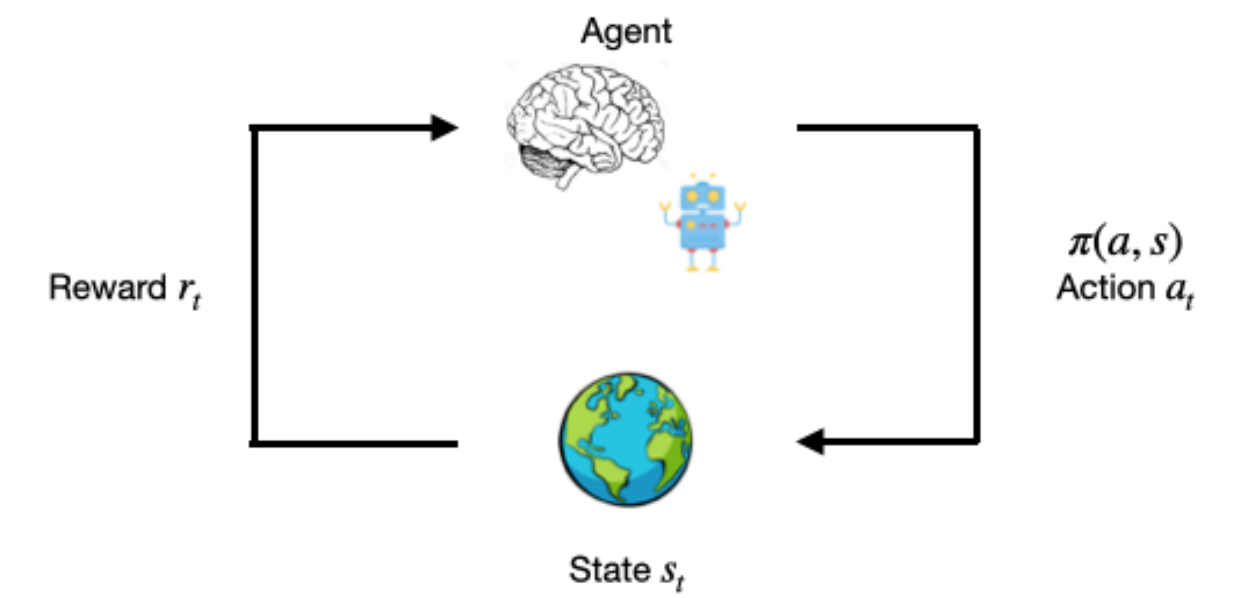


Fundamental property: **Markov property**


$$P(s_{t+1} = s \mid s_t, s_{t-1}, s_{t-2}, \dots) = P(s_{t+1} = s \mid s_t)$$

“The future is independent of the past given the present”




# Markov Process



Does this problem have the **Markov Property**?



$$P(s_{t+1} = s \mid s_t, s_{t-1}, s_{t-2}, \dots) \stackrel{???}{=} P(s_{t+1} = s \mid s_t)$$



???

# MDPs basis for model-based RL

$$P(s', r | s, a) = P(s_{t+1} = s', r_{t+1} = r | s_t = s, a_t = a)$$

How can we make use of such models of the world?

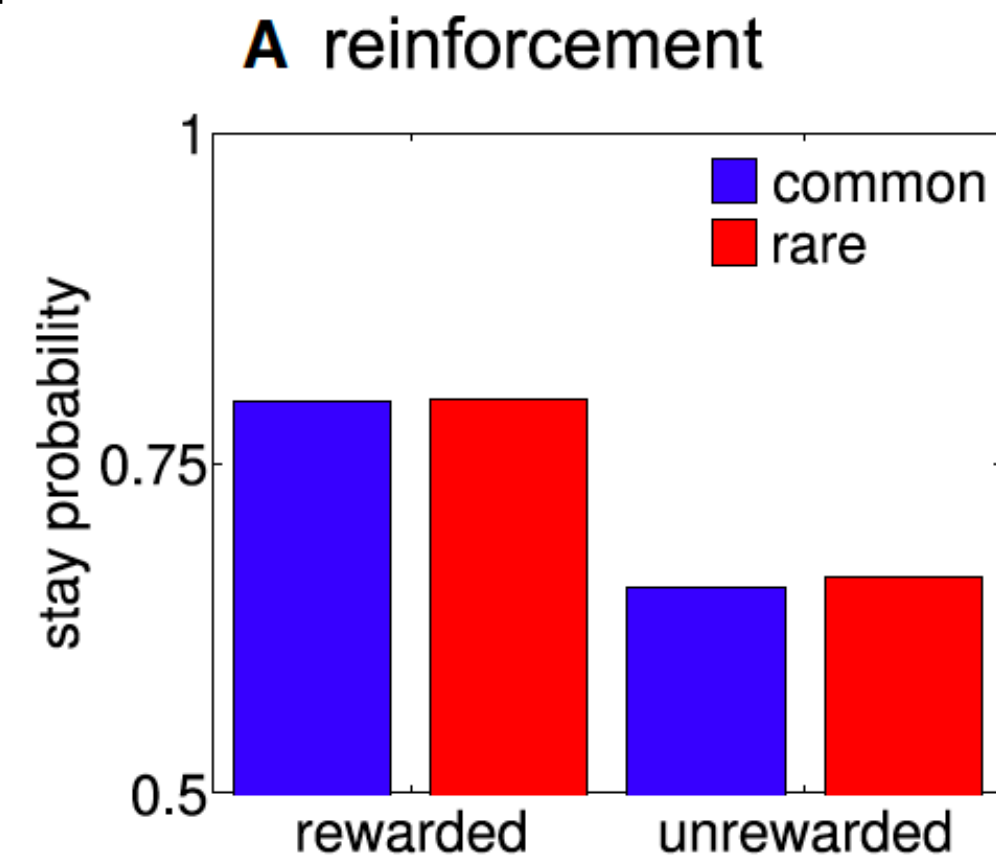


**Planning** and **action selection**

**Learning**

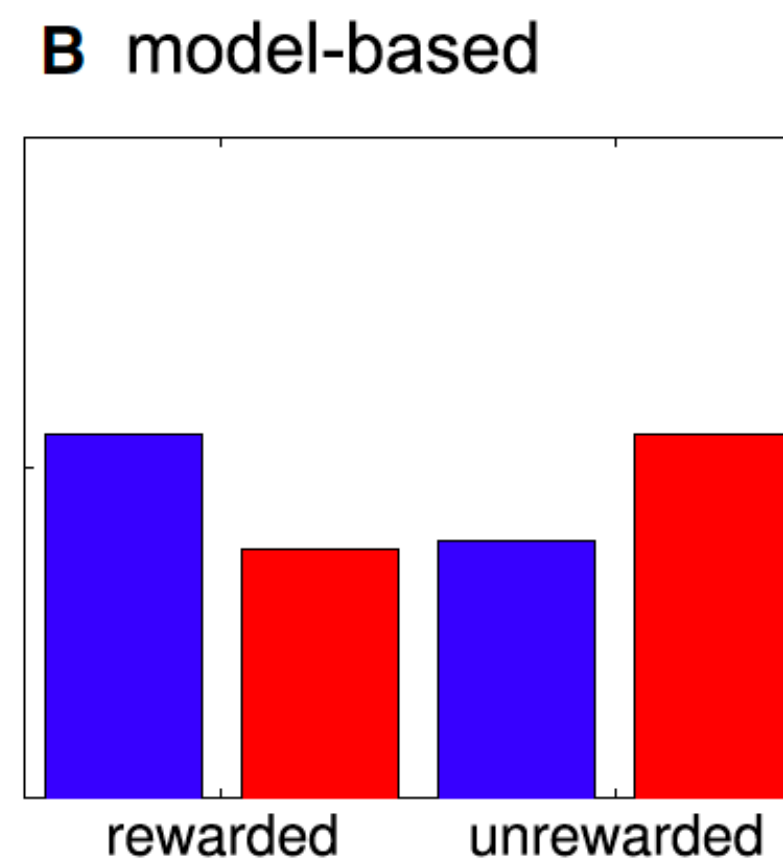
- Key idea: store experiences in world model  $P(s', r | s, a)$

# Two-step task: one of the most iconic RL tasks


Distinguish model-free vs. Model-based learning



Model-free RL agent: repeat what is rewarding



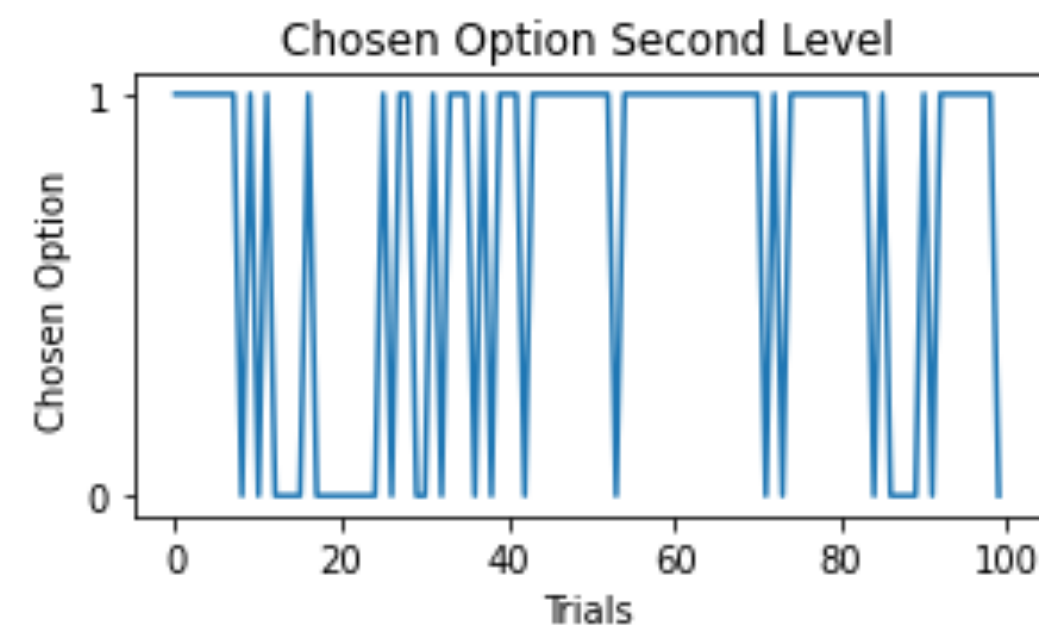
Model-based RL agent: repeat what is rewarding, but be clever




Really data: a mix of both

# Problem: how do we find the best parameters for a given model?


Assume your participant behaves like this:




(Here: data generated with  $\alpha = 0.5$  and  $\beta = 5$ )



( $\alpha = 0$  and  $\beta = 5$ )



( $\alpha = 1$  and  $\beta = 5$ )



( $\alpha = 0.55$  and  $\beta = 5.06$ )

***Any other Questions?***